

# DL Pac 1

```
In [1]: import numpy as np  
import pandas as pd
```

```
In [13]: from sklearn.datasets import fetch_openml  
  
# Load Boston housing dataset  
boston = fetch_openml(name="boston")  
  
# Converting the data into pandas DataFrame  
data= pd.DataFrame(boston.data, columns=boston.feature_names)  
  
# Adding the target variable to the dataset  
data['PRICE'] = boston.target  
  
# First Look at the data  
print(data.head())  
  
# Shape of the data  
print(data.shape)
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296.0	15.3	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242.0	17.8	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242.0	17.8	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222.0	18.7	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222.0	18.7	

	B	LSTAT	PRICE
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	5.33	36.2

(506, 14)

C:\Users\SOFT LAB11\anaconda3\Lib\site-packages\sklearn\datasets\\_openml.py:301: UserWarning: Multiple active versions of the dataset matching the name boston exist. Versions may be fundamentally different, returning version 1.

    warn(

C:\Users\SOFT LAB11\anaconda3\Lib\site-packages\sklearn\datasets\\_openml.py:968: FutureWarning: The default value of `parser` will change from ``liac-arff`` to ``auto`` in 1.4. You can set `parser='auto'` to silence this warning. Therefore, an `ImportError` will be raised from 1.4 if the dataset is dense and pandas is not installed. Note that the pandas parser may return different data types. See the Notes Section in fetch\_openml's API doc for details.

    warn(

```
In [14]: #Converting the data into pandas dataframe  
data= pd.DataFrame(boston.data)  
#First Look at the data  
data.head()
```

Out[14]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	296.0	15.3	396.90	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242.0	17.8	396.90
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242.0	17.8	392.83
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222.0	18.7	394.63
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222.0	18.7	396.90



In [15]: #Adding the feature names to the dataframe

```
data.columns= boston.feature_names  
#Adding the target variable to the dataset  
data['PRICE'] = boston.target  
#Looking at the data with names and target variable  
data.head(n=10)
```

Out[15]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	296.0	15.3	396.90	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242.0	17.8	396.90
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242.0	17.8	392.83
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222.0	18.7	394.63
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222.0	18.7	396.90
5	0.02985	0.0	2.18	0	0.458	6.430	58.7	6.0622	3	222.0	18.7	394.12
6	0.08829	12.5	7.87	0	0.524	6.012	66.6	5.5605	5	311.0	15.2	395.60
7	0.14455	12.5	7.87	0	0.524	6.172	96.1	5.9505	5	311.0	15.2	396.90
8	0.21124	12.5	7.87	0	0.524	5.631	100.0	6.0821	5	311.0	15.2	386.63
9	0.17004	12.5	7.87	0	0.524	6.004	85.9	6.5921	5	311.0	15.2	386.71



In [16]: #Shape of the data

```
print(data.shape)  
#Checking the null values in the dataset  
data.isnull().sum()
```

(506, 14)

```
Out[16]: CRIM      0  
ZN         0  
INDUS     0  
CHAS      0  
NOX       0  
RM         0  
AGE       0  
DIS        0  
RAD        0  
TAX        0  
PTRATIO   0  
B          0  
LSTAT     0  
PRICE     0  
dtype: int64
```

```
In [17]: #Checking the statistics of the data  
data.describe()
```

```
Out[17]:
```

	CRIM	ZN	INDUS	NOX	RM	AGE	DIS
<b>count</b>	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
<b>mean</b>	3.613524	11.363636	11.136779	0.554695	6.284634	68.574901	3.795043
<b>std</b>	8.601545	23.322453	6.860353	0.115878	0.702617	28.148861	2.105710
<b>min</b>	0.006320	0.000000	0.460000	0.385000	3.561000	2.900000	1.129600
<b>25%</b>	0.082045	0.000000	5.190000	0.449000	5.885500	45.025000	2.100175
<b>50%</b>	0.256510	0.000000	9.690000	0.538000	6.208500	77.500000	3.207450
<b>75%</b>	3.677083	12.500000	18.100000	0.624000	6.623500	94.075000	5.188425
<b>max</b>	88.976200	100.000000	27.740000	0.871000	8.780000	100.000000	12.126500

```
In [18]: data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
 #   Column  Non-Null Count Dtype  
---- 
 0   CRIM    506 non-null   float64 
 1   ZN      506 non-null   float64 
 2   INDUS   506 non-null   float64 
 3   CHAS    506 non-null   category 
 4   NOX     506 non-null   float64 
 5   RM      506 non-null   float64 
 6   AGE     506 non-null   float64 
 7   DIS     506 non-null   float64 
 8   RAD     506 non-null   category 
 9   TAX     506 non-null   float64 
 10  PTRATIO 506 non-null   float64 
 11  B       506 non-null   float64 
 12  LSTAT   506 non-null   float64 
 13  PRICE   506 non-null   float64 
dtypes: category(2), float64(12)
memory usage: 49.0 KB

```

In [19]: pip install seaborn

```

Requirement already satisfied: seaborn in c:\users\soft lab1\anaconda3\lib\site-packages (0.12.2)
Requirement already satisfied: numpy!=1.24.0,>=1.17 in c:\users\soft lab1\anaconda3\lib\site-packages (from seaborn) (1.26.4)
Requirement already satisfied: pandas>=0.25 in c:\users\soft lab1\anaconda3\lib\site-packages (from seaborn) (2.1.4)
Requirement already satisfied: matplotlib!=3.6.1,>=3.1 in c:\users\soft lab1\anaconda3\lib\site-packages (from seaborn) (3.8.0)
Requirement already satisfied: contourpy>=1.0.1 in c:\users\soft lab1\anaconda3\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (1.2.0)
Requirement already satisfied: cycler>=0.10 in c:\users\soft lab1\anaconda3\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (0.11.0)
Requirement already satisfied: fonttools>=4.22.0 in c:\users\soft lab1\anaconda3\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (4.25.0)
Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\soft lab1\anaconda3\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (1.4.4)
Requirement already satisfied: packaging>=20.0 in c:\users\soft lab1\anaconda3\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (23.1)
Requirement already satisfied: pillow>=6.2.0 in c:\users\soft lab1\anaconda3\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (10.2.0)
Requirement already satisfied: pyparsing>=2.3.1 in c:\users\soft lab1\anaconda3\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (3.0.9)
Requirement already satisfied: python-dateutil>=2.7 in c:\users\soft lab1\anaconda3\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in c:\users\soft lab1\anaconda3\lib\site-packages (from pandas>=0.25->seaborn) (2023.3.post1)
Requirement already satisfied: tzdata>=2022.1 in c:\users\soft lab1\anaconda3\lib\site-packages (from pandas>=0.25->seaborn) (2023.3)
Requirement already satisfied: six>=1.5 in c:\users\soft lab1\anaconda3\lib\site-packages (from python-dateutil>=2.7->matplotlib!=3.6.1,>=3.1->seaborn) (1.16.0)
Note: you may need to restart the kernel to use updated packages.

```

```
In [20]: #checking the distribution of the target variable
```

```
import seaborn as sns  
sns.distplot(data.PRICE)
```

```
C:\Users\SOFT LAB11\AppData\Local\Temp\ipykernel_8924\4212025153.py:3: UserWarning:
```

```
'distplot' is a deprecated function and will be removed in seaborn v0.14.0.
```

```
Please adapt your code to use either 'displot' (a figure-level function with similar flexibility) or 'histplot' (an axes-level function for histograms).
```

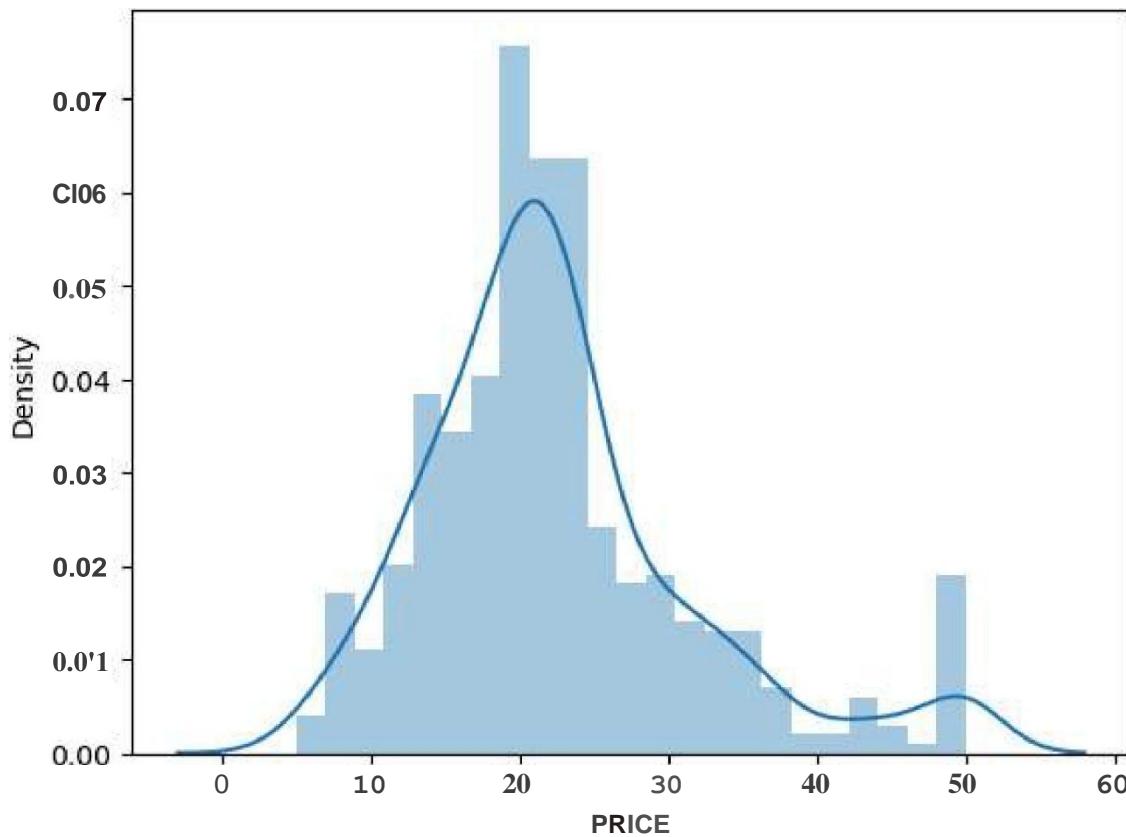
```
For a guide to updating your code to use the new functions, please see  
https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751
```

```
sns.distplot(data.PRICE)
```

```
C:\Users\SOFT LAB11\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
```

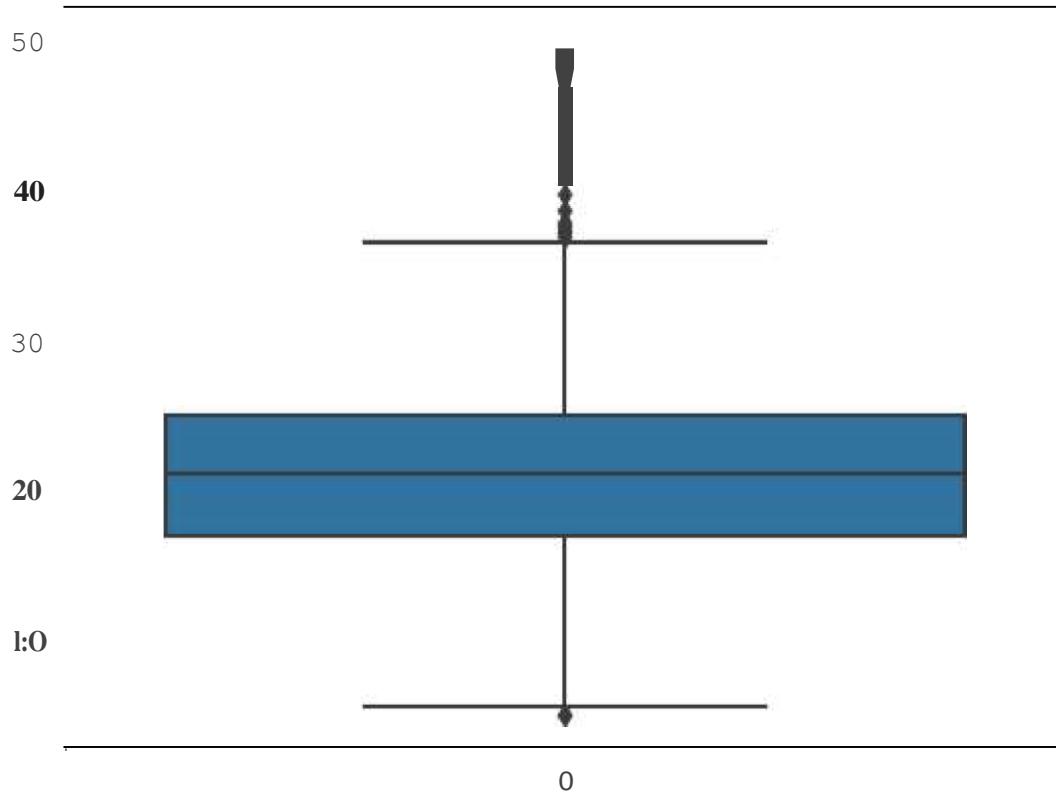
```
with pd.option_context('mode.use_inf_as_na', True):
```

```
Out[20]: <Axes: xlabel='PRICE', ylabel='Density'>
```



```
In [21]: sns.boxplot(data.PRICE)
```

```
Out[21]: <Axes: >
```



```
In [22]: correlation = data.corr()  
correlation.loc['PRICE']
```

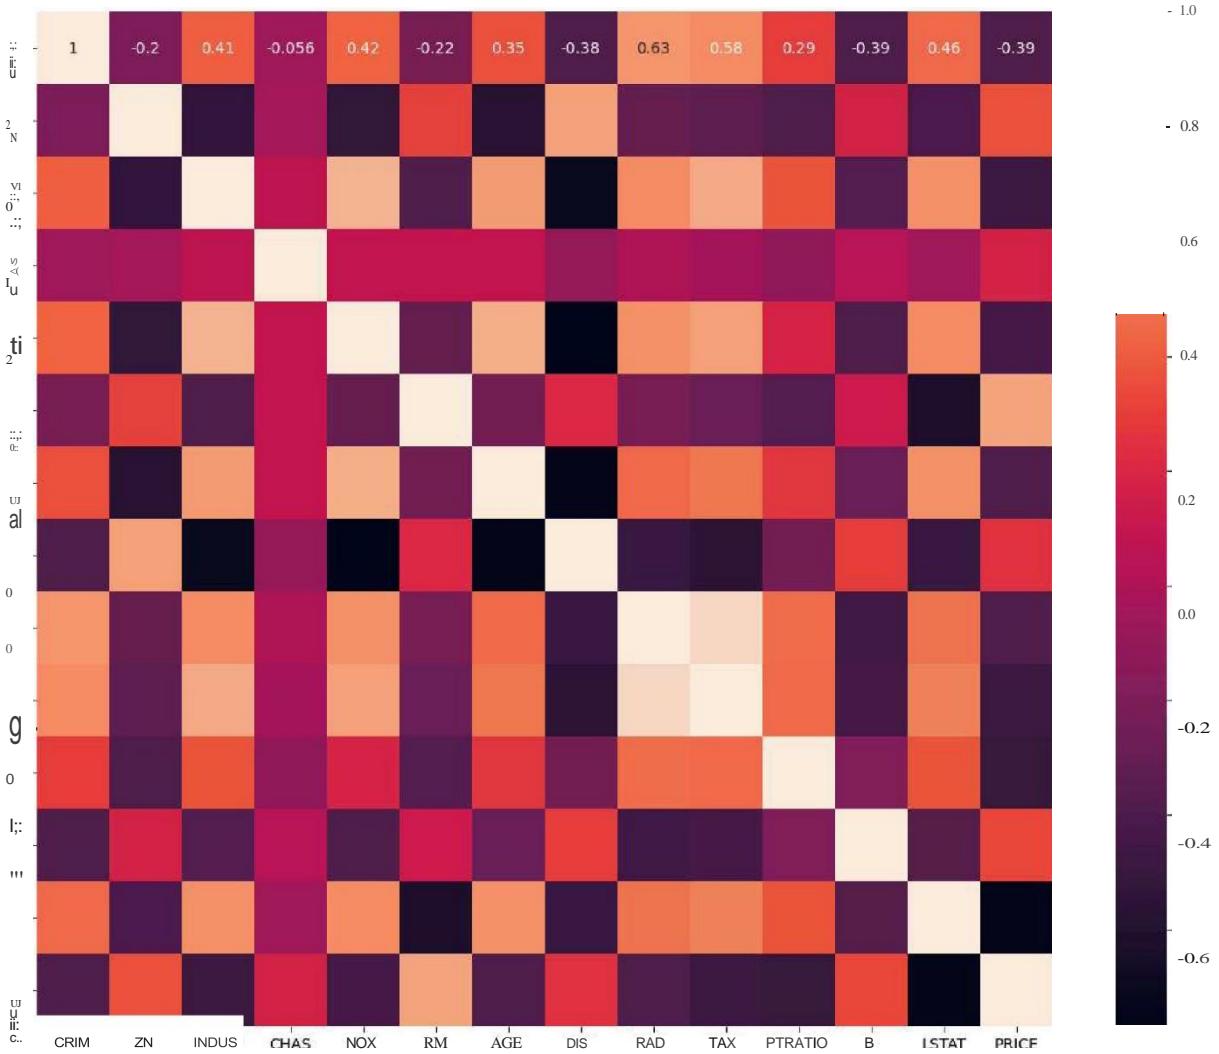
```
Out[22]: CRIM      -0.388305  
ZN         0.360445  
INDUS     -0.483725  
CHAS       0.175260  
NOX        -0.427321  
RM         0.695360  
AGE        -0.376955  
DIS        0.249929  
RAD        -0.381626  
TAX        -0.468536  
PTRATIO   -0.507787  
B          0.333461  
LSTAT    -0.737663  
PRICE      1.000000
```

Name: PRICE, dtype: float64

```
In [23]: # plotting the heatmap  
import matplotlib.pyplot as plt
```

```
In [24]: fig,axes = plt.subplots(figsize=(15,12))  
sns.heatmap(correlation,square = True,annot = True)
```

```
Out[24]: <Axes: >
```



```
In [25]: import matplotlib.pyplot as plt
```

```
# Checking the scatter plot with the most correlated features
plt.figure(figsize=(20, 5))
features = ['LSTAT', 'RM', 'PTRATIO']

for i, col in enumerate(features):
    plt.subplot(1, len(features), i + 1)
    x = data[col]
    y = data.PRICE
    plt.scatter(x, y, marker='o')
    plt.title("Variation in House prices")
    plt.xlabel(col)
    plt.ylabel("House prices in $1000")

plt.show()
```



```
In [26]: # Splitting the dependent feature and independent feature
#X = data[['LSTAT','RM', 'PTRATIO']]
```

```
from sklearn.model_selection import train_test_split

X = data.iloc[:, :-1]
y = data.PRICE
```

```
In [27]: # Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [28]: # Assuming X_train and X_test are DataFrames

# Exclude non-numeric columns from standardization
numeric_columns = X_train.select_dtypes(include=[np.number]).columns

# Calculate mean and standard deviation only for numeric columns
mean = X_train[numeric_columns].mean(axis=0)
std = X_train[numeric_columns].std(axis=0)

# Standardize the training data for numeric columns
X_train_standardized = (X_train[numeric_columns] - mean) / std

# Standardize the testing data using the mean and std from the training data for numeric columns
X_test_standardized = (X_test[numeric_columns] - mean) / std
```

```
In [46]: # Assuming X and y are already defined and Loaded

# Exclude non-numeric columns from standardization
numeric_columns = X.select_dtypes(include=[np.number]).columns

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X[numeric_columns], y, test_size=0.2, random_state=42)

# Standardize the features
mean = X_train.mean(axis=0)
std = X_train.std(axis=0)
X_train = (X_train - mean) / std
X_test = (X_test - mean) / std

# Linear Regression
regressor = LinearRegression()

# Fitting the model
regressor.fit(X_train, y_train)
```

```
regressor.fit(X_train, y_train)

# Model Evaluation
# Prediction on the test dataset
y_pred = regressor.predict(X_test)

# Calculate RMSE (Root Mean Squared Error)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
print("RMSE:", rmse)

# Calculate R-squared (R2) score
r2 = r2_score(y_test, y_pred)
print("R-squared (R2) Score:", r2)

RMSE: 5.203503176683114
R-squared (R2) Score: 0.6307780105854284
```

```
In [29]: # Neural Networks
#Scaling the dataset
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

```
In [30]: !pip install graphviz
!pip install ann_visualizer

Collecting graphviz
  Downloading graphviz-0.20.3-py3-none-any.whl.metadata (12 kB)
  Downloading graphviz-0.20.3-py3-none-any.whl (47 kB)
----- 0.0/47.1 kB? eta:-----
----- 20.5/47.1 kB 682.7 kB/s eta 0:00:01
----- 47.1/47.1 kB 594.9 kB/s eta 0:00:00
Installing collected packages: graphviz
Successfully installed graphviz-0.20.3
Collecting ann_visualizer

  Downloading ann_visualizer-2.5.tar.gz (4.7 kB)
  Preparing metadata (setup.py): started
  Preparing metadata (setup.py): finished with status 'done'
Building wheels for collected packages: ann_visualizer
  Building wheel for ann_visualizer (setup.py): started
  Building wheel for ann_visualizer (setup.py): finished with status 'done'
  Created wheel for ann_visualizer: filename=ann_visualizer-2.5-py3-none-any.whl size=4181 sha256=150bbe984af6cfbf7807a5f643f6232f511556d1fc33a3f41af8c46d79ef32
  Stored in directory: c:\users\soft lab11\appdata\local\pip\cache\wheels\28\4a\ad\82dalaad2994e42bf0f4bld403fd8a64dfc38ae2c8a5daa4
Successfully built ann_visualizer
Installing collected packages: ann_visualizer
Successfully installed ann_visualizer-2.5
```

```
In [32]: pip install keras
```

```
In [34]: pip install tensorflow
```

```
In [35]: from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import Dense
```

```
In [36]: import keras  
from keras.models import Sequential  
from keras.layers import Dense
```

```
In [37]: # Importing necessary libraries  
import keras  
from keras.layers import Dense  
from keras.models import Sequential  
from keras.optimizers import Adam  
from keras.callbacks import EarlyStopping  
from ann_visualizer.visualize import ann_viz  
import plotly.subplots as sp  
import plotly.graph_objects as go  
  
# Creating the neural network model  
model = Sequential()  
model.add(Dense(128, activation='relu', input_dim=13))  
model.add(Dense(64, activation='relu'))  
model.add(Dense(32, activation='relu'))  
model.add(Dense(16, activation='relu'))  
model.add(Dense(1))  
  
# Compiling the model  
model.compile(optimizer='adam', loss='mean_squared_error', metrics=['mae'])
```

```
C:\Users\SOFT LAB11\anaconda3\Lib\site-packages\keras\src\layers\core\dense.py:87: UserWarning: Do not pass an 'input_shape'/'input_dim' argument to a layer. When using Sequential models, prefer using an 'Input(shape)' object as the first layer in them odel instead.
```

```
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
In [38]: from ann_visualizer.visualize import ann_viz
```

```
In [39]: pip install keras matplotlib plotly
```

```
In [40]: import numpy as np  
import pandas as pd  
from sklearn.datasets import fetch_openml  
from sklearn.model_selection import train_test_split  
from sklearn.preprocessing import StandardScaler  
from tensorflow.keras.models import Sequential
```

```

from tensorflow.keras.layers import Dense
from tensorflow.keras.callbacks import EarlyStopping
import plotly.graph_objects as go

# Load Boston housing dataset
boston = fetch_openml(name="boston")

# Converting the data into pandas DataFrame
data= pd.DataFrame(boston.data, columns=boston.feature_names)

# Adding the target variable to the dataset
data['PRICE'] = boston.target

# Extract features (X) and target variable (y)
X = data.drop(columns=['PRICE']).values
y = data['PRICE'].values

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardize the features
scaler= StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Define the neural network model
model= Sequential([
    Dense(64, activation='relu', input_shape=(X_train_scaled.shape[1],)),
    Dense(32, activation='relu'),
    Dense(1) # Output Layer with single neuron for regression
])

# Compile the model
model.compile(optimizer='adam', loss='mean_squared_error', metrics=['mae'])

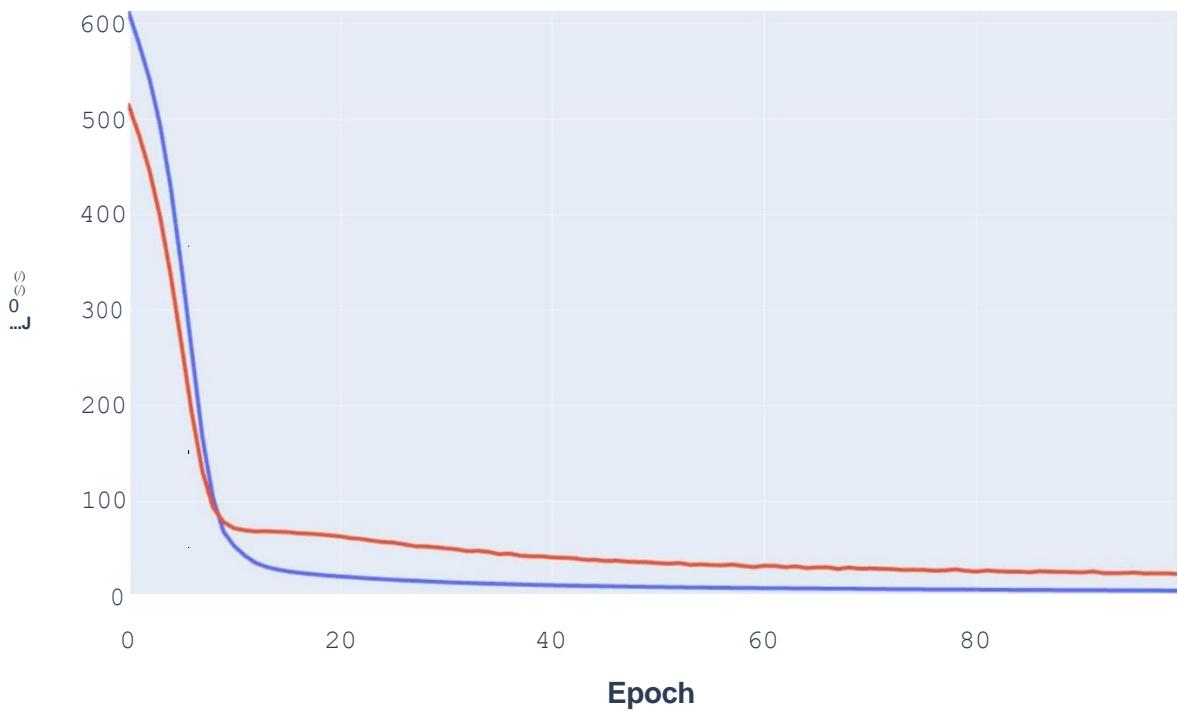
# Training the model
history= model.fit(X_train_scaled, y_train, epochs=100, validation_split=0.05, callbacks=[EarlyStopping(monitor='val_mae', patience=10, restore_best_weights=True)])

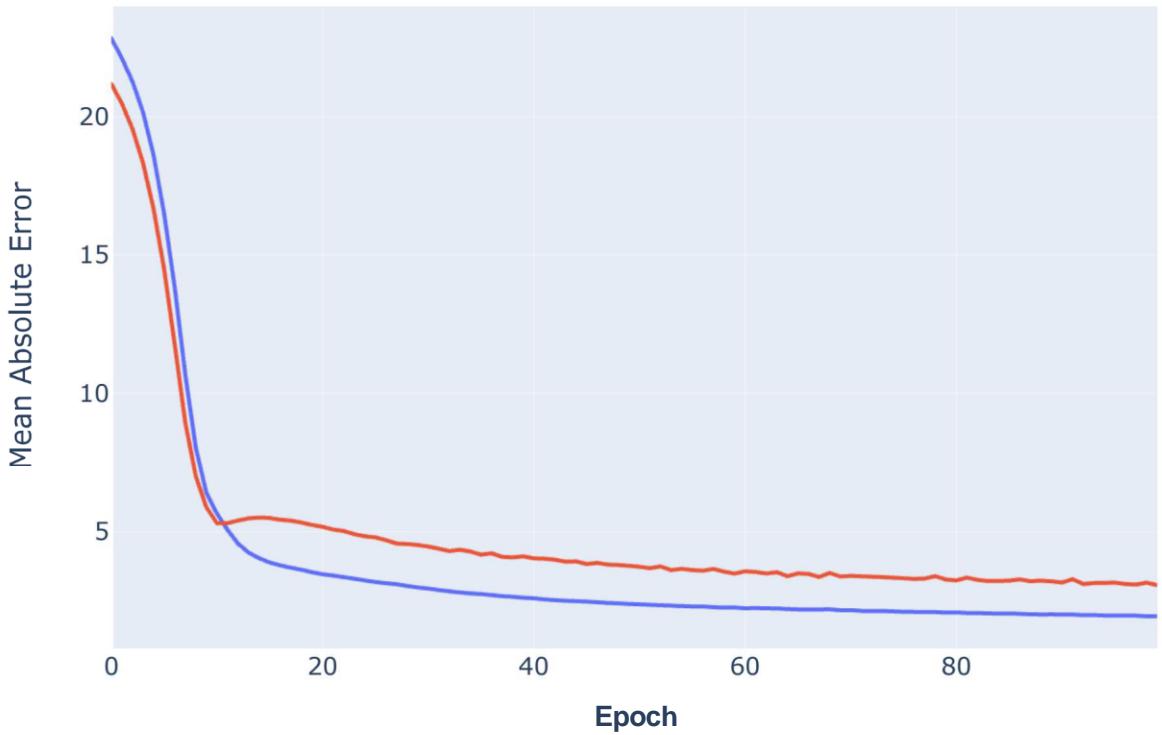
# Plotting the training and validation Loss
fig_loss = go.Figure()
fig_loss.add_trace(go.Scattergl(y=history.history['loss'], name='Train'))
fig_loss.add_trace(go.Scattergl(y=history.history['val_loss'], name='Valid'))
fig_loss.update_layout(height=500, width=700, xaxis_title='Epoch', yaxis_title='Loss')
fig_loss.show()

# Plotting the training and validation mean absolute error
fig_mae = go.Figure()
fig_mae.add_trace(go.Scattergl(y=history.history['mae'], name='Train'))
fig_mae.add_trace(go.Scattergl(y=history.history['val_mae'], name='Valid'))
fig_mae.update_layout(height=500, width=700, xaxis_title='Epoch', yaxis_title='Mean Absolute Error')
fig_mae.show()

```

Epoch 1/100





◀ ▶

```
In [44]: # Convert y_test to a numpy array if it's not already
y_test = np.array(y_test)

# Evaluate the model on the test data
mse_nn, mae_nn = model.evaluate(X_test_scaled, y_test)

# Print the evaluation metrics
print('Mean squared error on test data: ', mse_nn)
print('Mean absolute error on test data: ', mae_nn)
```

**4/4** ————— 0s 7ms/step - loss: 9.2586 - mae: 2.2224

Mean squared error on test data: 12.195608139038086

Mean absolute error on test data: 2.3551769256591797

```
In [45]: import numpy as np
import pandas as pd
from sklearn.datasets import fetch_openml
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.callbacks import EarlyStopping
import plotly.graph_objects as go
```

```

# Load Boston housing dataset
boston = fetch_openml(name="boston")

# Converting the data into pandas DataFrame
data= pd.DataFrame(boston.data, columns=boston.feature_names)

# Adding the target variable to the dataset
data['PRICE'] = boston.target

# Extract features (X) and target variable (y)
X = data.drop(columns=['PRICE']).values
y = data['PRICE'].values

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_stan

# Standardize the features
scaler= StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Define the neural network model
model= Sequential([
    Dense(64, activation='relu', input_shape=(X_train_scaled.shape[1],)),
    Dense(32, activation='relu'),
    Dense(1) # Output Layer with single neuron for regression
])

# Compile the model
model.compile(optimizer='adam', loss='mean_squared_error', metrics=['mae'])

# Training the model
history= model.fit(X_train_scaled, y_train, epochs=100, validation_split=0.05, cal

# Evaluation of the neural network model
y_pred_nn = model.predict(X_test_scaled)
mse_nn = mean_squared_error(y_test, y_pred_nn)
mae_nn = mean_absolute_error(y_test, y_pred_nn)
r2_nn = r2_score(y_test, y_pred_nn)

print('Neural Network Model:')
print('Mean squared error on test data:', mse_nn)
print('Mean absolute error on test data:', mae_nn)
print('R-squared (R2) score:', r2_nn)

# Comparison with Linear Regression
lr_model = LinearRegression()
lr_model.fit(X_train_scaled, y_train)
y_pred_lr = lr_model.predict(X_test_scaled)
mse_lr = mean_squared_error(y_test, y_pred_lr)
mae_lr = mean_absolute_error(y_test, y_pred_lr)
r2_lr = r2_score(y_test, y_pred_lr)

print('\nLinear Regression Model:')
print('Mean squared error on test data:', mse_lr)

```

```
print('Mean absolute error on test data:', mae_lr)
print('R-squared (R2) score:', r2_lr)
```

In [46]: # Make predictions on new data

```
import sklearn
new_data = sklearn.preprocessing.StandardScaler().fit_transform([[0.1, 10.0,
5.0, 0, 0.4, 6.0, 50, 6.0, 1, 400, 20, 300, 10]]))
prediction= model.predict(new_data)
print("Predicted house price:", prediction)
```

1/1  3ms/step

-----  
Predicted house price: [[9.126371]]

In [ ]: